

Project Tungsten Phase II

Joining a Billion Rows per Second on a Laptop

Sameer Agarwal

Spark Meetup | SAP | June 30th 2016

About Me

- Software Engineer at Databricks (Spark Core/SQL)
- PhD in Databases (UC Berkeley)
- Research on BlinkDB (Approximate Queries in Spark)

Hardware Trends

Storage

Network

CPU

Hardware Trends

2010

Storage 50+MB/s
(HDD)

Network 1Gbps

CPU ~3GHz

Hardware Trends

	2010	2016
Storage	50+MB/s (HDD)	500+MB/s (SSD)
Network	1Gbps	10Gbps
CPU	~3GHz	~3GHz

Hardware Trends

	2010	2016	
Storage	50+MB/s (HDD)	500+MB/s (SSD)	10X
Network	1Gbps	10Gbps	10X
CPU	~3GHz	~3GHz	😞

On the flip side

Spark IO has been optimized

- Reduce IO by pruning input data that is not needed
- New shuffle and network implementations (2014 sort record)


Data formats have improved

- E.g. Parquet is a “dense” columnar format

CPU increasingly the bottleneck; trend expected to continue

Goals of Project Tungsten

Substantially improve the **memory and CPU** efficiency of Spark **backend execution** and push performance closer to the limits of modern hardware.



Note the focus on “execution” not “optimizer”: very easy to pick broadcast join that is 1000X faster than Cartesian join, but hard to optimize broadcast join to be an order of magnitude faster.

Phase 1 Foundation

Memory Management
Code Generation
Cache-aware Algorithms

Phase 2 Order-of-magnitude Faster

Whole-stage Codegen
Vectorization

The background is a textured teal watercolor wash. It features darker, more saturated teal areas in the upper half, which blend into lighter, more translucent teal towards the bottom. The texture is organic and painterly, with visible brushstrokes and color variations.

Phase 1

Laying The Foundation

Summary

Perform manual memory management instead of relying on Java objects

- Reduce memory footprint
- Eliminate garbage collection overheads
- Use `java.unsafe` and off heap memory

Code generation for expression evaluation

- Reduce virtual function calls and interpretation overhead

Cache conscious sorting

- Reduce bad memory access patterns

The background of the slide is a textured teal or turquoise color, resembling a watercolor wash. The color is not uniform, with darker, more saturated areas in the upper left and center, and lighter, more translucent areas towards the bottom and right. The texture is soft and painterly, with visible brushstrokes and color blending.

Phase 2

Order-of-magnitude Faster


Going back to the fundamentals

Difficult to get order of magnitude performance speed ups with profiling techniques

- For 10x improvement, would need of find top hotspots that add up to 90% and make them instantaneous
- For 100x, 99%

Instead, look bottom up, how fast should *it* run?

```
select count(*) from store_sales  
where ss_item_sk = 1000
```



Volcano Iterator Model

Standard for 30 years: almost all databases do it

Each operator is an “iterator”
that consumes records from
its input operator

```
class Filter(  
    child: Operator,  
    predicate: (Row => Boolean))  
    extends Operator {  
    def next(): Row = {  
        var current = child.next()  
        while (current == null || predicate(current)) {  
            current = child.next()  
        }  
        return current  
    }  
}
```

Downside of the Volcano Model

1. Too many virtual function calls
 - at least 3 calls for each row in Aggregate
2. Extensive memory access
 - “row” is a small segment in memory (or in L1/L2/L3 cache)
3. Can't take advantage of modern CPU features
 - SIMD, pipelining, prefetching, branch prediction, ILP, instruction cache, ...

What if we hire a college freshman to implement this query in Java in 10 mins?


```
select count(*) from store_sales  
where ss_item_sk = 1000
```

```
long count = 0;  
for (ss_item_sk in store_sales) {  
    if (ss_item_sk == 1000) {  
        count += 1;  
    }  
}
```

Volcano model
30+ years of database research

vs

college freshman
hand-written code in 10 mins



How does a student beat 30 years of research?

Volcano

1. Many virtual function calls
2. Data in memory (or cache)
3. No loop unrolling, SIMD, pipelining

hand-written code

1. No virtual function calls
2. Data in CPU registers
3. Compiler loop unrolling, SIMD, pipelining

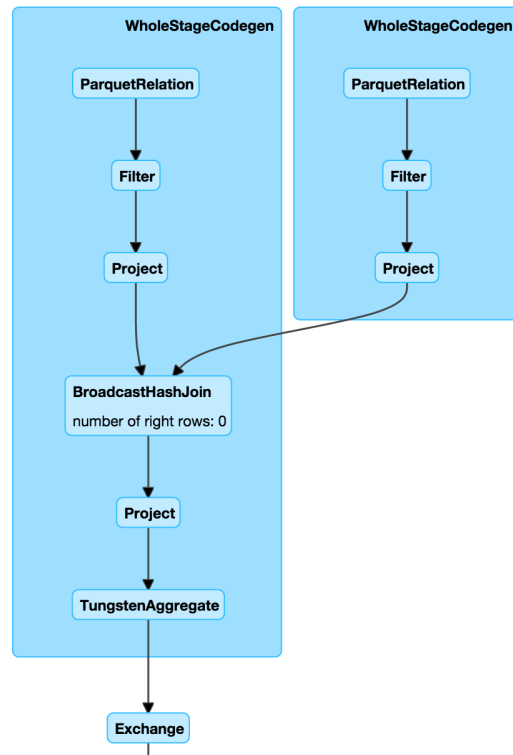
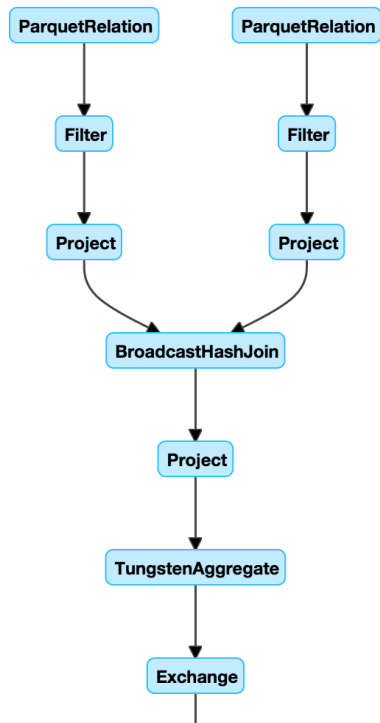
Take advantage of all the information that is known **after** query compilation

Whole-stage Codegen

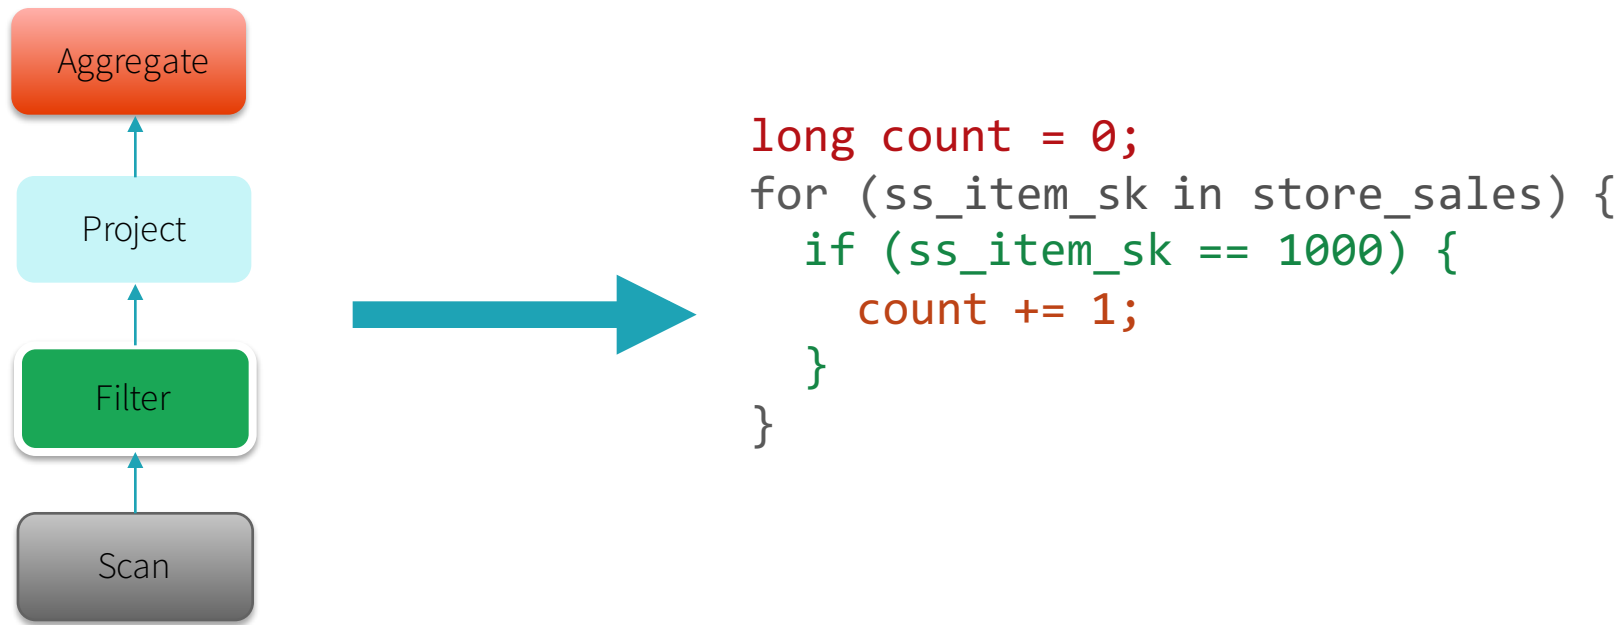
Fusing operators together so the generated code looks like hand optimized code:

- Identify chains of operators (“stages”)
- Compile each stage into a single function
- Functionality of a general purpose execution engine; performance as if hand built system just to run your query

Whole-stage Codegen: Planner



Whole-stage Codegen: Spark as a “Compiler”



But there are things we can't fuse

Complicated I/O

- CSV, Parquet, ORC, ...
- Sending across the network

External integrations

- Python, R, scikit-learn, TensorFlow, etc
- Reading cached data

Columnar in memory format

In-memory
Row Format

1	john	4.1
2	mike	3.5
3	sally	6.4

In-memory
Column Format

1	2	3
john	mike	sally
4.1	3.5	6.4

Why columnar?

1. More efficient: denser storage, regular data access, easier to index into. Enables vectorized processing.
2. More compatible: Most high-performance external systems are already columnar (numpy, TensorFlow, Parquet); zero serialization/copy to work with them
3. Easier to extend: process encoded data, integrate with columnar cache etc.



The background is a textured teal watercolor wash. It features darker, more saturated teal areas in the upper half, which blend into lighter, more transparent teal towards the bottom. The texture is organic and painterly, with visible brushstrokes and color variations.

Putting it *All* Together

Phase 1
Spark 1.4 - 1.6

Memory Management
Code Generation
Cache-aware Algorithms


Phase 2
Spark 2.0+

Whole-stage Code Generation
Columnar in Memory Support

Operator Benchmarks: Cost/Row (ns)

primitive	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8-bit entropy)	620 ns	5.3 ns
sort (64-bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns
Parquet decoding (single int column)	120 ns	13 ns

TPC-DS (Scale Factor 1500, 100 cores)



Status

- Being released as part of Spark 2.0
 - Both Whole stage codegen and vectorized Parquet reader is on by default
- Back to profiling techniques
 - Improve quality of generated code, optimize Parquet reader further
- Try it out and let us know!

Further Reading

Apache Spark as a Compiler: Joining a Billion Rows per Second on a Laptop

Deep dive into the new Tungsten execution engine



by Sameer Agarwal, Davies Liu and Reynold Xin

Posted in **ENGINEERING BLOG** | May 23, 2016

Spark Summit 2016 will be held in San Francisco on June 6–8. Check out the [full agenda](#) and [get your ticket](#) before it sells out!

2016 Apache Spark Survey



Survey 2016

Share your thoughts in our 2016
Apache® Spark™ Survey.

TAKE THE SURVEY

Spark Summit EU
Brussels
October 25-27

The CFP closes at **11:59pm on July 1st**
For more information and to submit:

<https://spark-summit.org/eu-2016/>

Questions?